



Team id : BD_1482_1903_1957_1972

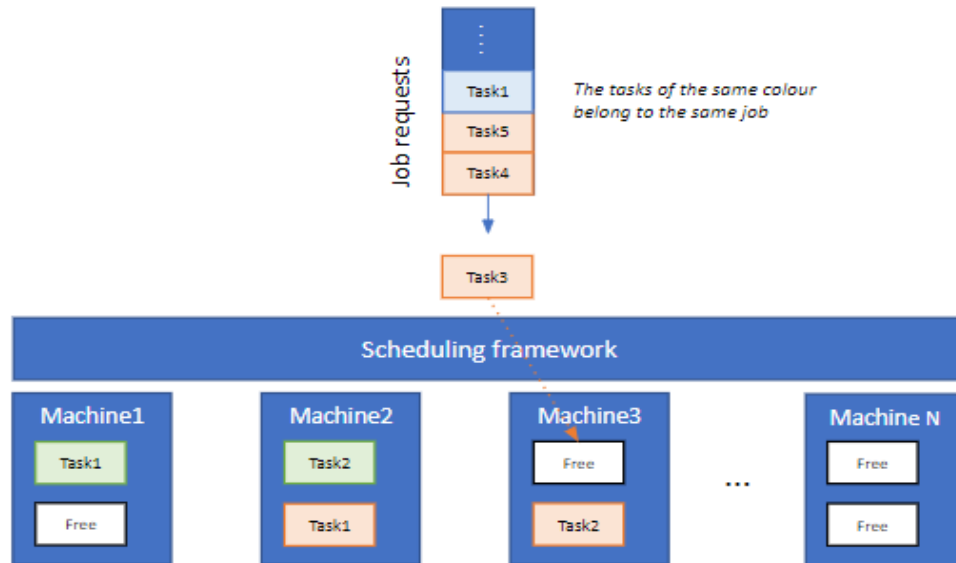
Final Class Project Report

Date:01/12/2020

SNo	Name	SRN	Class/Section
1	S Nikhil Ram	PES1201801972	5 -B
2	Rakshith C	PES1201801903	5 - D

3	K Vikas Gowda	PES1201801957	5 - F
4	Yash Gawankar	PES1201801482	5 - J

Introduction



- Big Data workloads consist of multiple jobs from different applications, which need to be run on a cluster of interconnected machines known as workers due to its large resource requirements.
- In the Real World, due to the limitation of resources, a scheduling framework is necessary in order to allocate resources to the cluster to various incoming jobs arriving from a wide range of applications.
- An incoming job could contain multiple tasks (like Map and Reduce), and hence resources need to be allocated in compliance with their dependencies (ex. all Reduce tasks depend on the previous Map tasks).
- Thus, in order to coordinate the scheduling of tasks among numerous interconnected clusters, a Centralized Scheduling Framework is needed.
- This framework consists of one **Master**, which runs on a dedicated machine and manages the resources of the rest of the machines in the cluster.
- The other machines in the cluster have one **Worker process** running on each of them .
- The Master process makes scheduling decisions while the Worker processes execute the tasks and inform the Master when a task completes its execution.

Related work

Research Papers:-

- Dazhao Cheng, Xiaobo Zhou, Palden Lama, Jun Wu, Changjun Jiang, Cross-platform resource scheduling for spark and MapReduce on YARN, IEEE Trans. Comput., PP (99) (2017)
- Jyoti V. Gautam, Harshad kumar B. Prajapati, Vipul K. Dabhi, Sanjay Chaudhary, A survey on job scheduling algorithms in big data processing, IEEE Conf. Pap. (March 2015)
- Thomas C. Bressoud, Qiuyi Tang, “Results of a model for hadoop YARN MapReduce tasks, IEEE Int. Conf. Clust. Comput. (2016)
- Mohd Usama, Mengchen Liu, Min Chen, Job Schedulers for Big Data Processing in Hadoop Environment: Testing real-life schedulers using benchmark programs
- Islam, Muhammed Tawfiqul & Buyya, Rajkumar, Resource Management and Scheduling for Big Data Applications in Cloud Computing Environments (2018).

Articles:-

- [Job Scheduling in Hadoop](#) - Medium.com Article
- [Hadoop - Schedulers and Type of Schedulers](#) - GeeksforGeeks
- [Schedulers in YARN : from concepts to configurations](#) - towardsdatascience.com
- [Hadoop Schedulers Tutorial - Job Scheduling in Hadoop](#) - data-flair.training

Books:-

- [Hadoop: The Definitive Guide, 4th Edition](#) - O'Reilly Publications, Chapter 4: YARN

Design

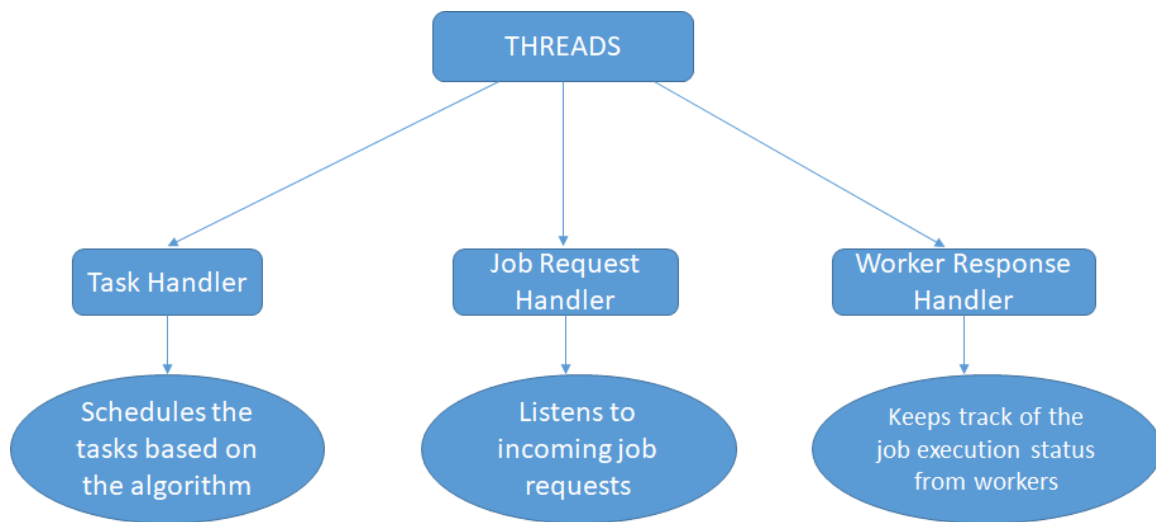
Overview of the General Design Strategy:-

- The Master listens for job requests and dispatches the tasks in the jobs to the workers based on a *scheduling algorithm*.
- The Worker processes listen for Task Launch messages from the Master.
- On receiving a Launch message, the Worker adds the task to the execution pool of the machine it runs on.
- The execution pool consists of all currently running tasks in the machine.

- When a task completes execution, the Worker process on the machine informs the Master.
- The Master then updates its information about the number of free slots available on the machine.

Implementation:-

Master :-



It is capable of performing 3 tasks simultaneously using threads:

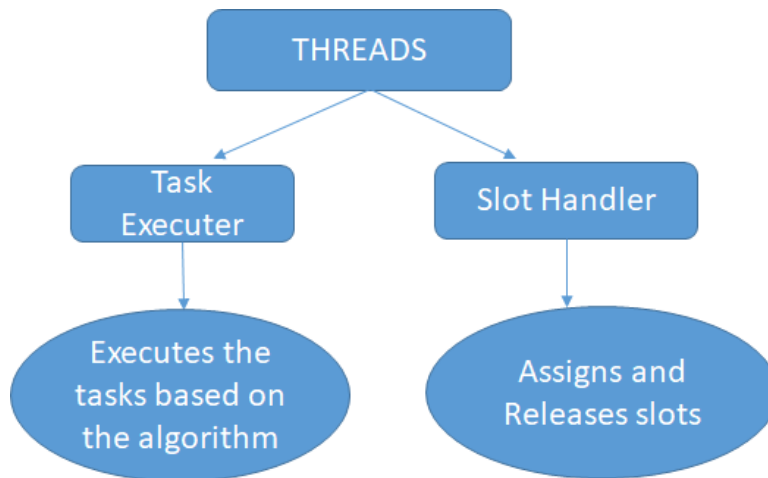
1. Listening to incoming Map and Reduce requests from the requests.py
2. Scheduling the tasks based on specified algorithms
3. Keeping track of execution status of the tasks and availability of resources on the workers

The threads are utilized by 3 classes namely:-

1. task_handler :
 - a. Connects to the worker based on specified port no.
 - b. Determines the type of algorithm used for scheduling
 - c. Schedules the jobs based on the Algorithm
 - d. It logs the start of each task in a job
 - e. Also keeps track of no of free slots on the workers
2. job_request_handler:

- a. Accepts the incoming job requests from requests.py
 - b. Reads and stores the JSON object consisting of Worker configurations
 - c. Appends the jobs to a job pool (queue)
 - d. Also logs the Start of every job
3. worker_response_handler:
 - a. Receives the task completion status from the workers
 - b. Logs the completion of each task
 - c. As and when the jobs are completed, it pops them from the job pool
 - d. Ensures that all Reduce tasks are implemented only after the Map tasks and thus maintaining the dependencies

Worker:-



It is capable of performing 2 tasks simultaneously using threads:

1. Accepting incoming Map and Reduce tasks from the master
2. Assigning available slots to each tasks

The threads are utilized by 2 classes namely:-

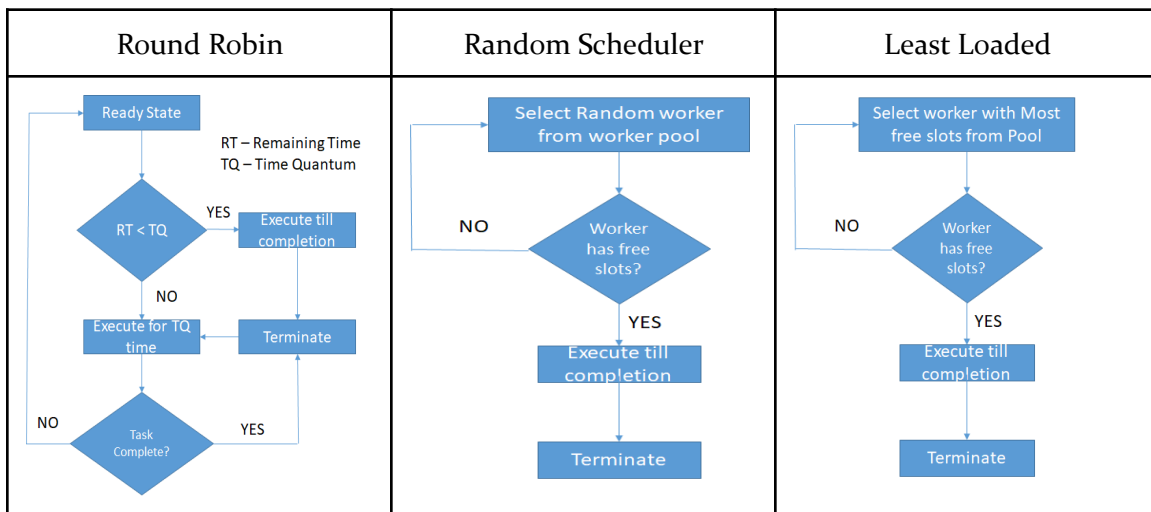
1. task_executer :
 - a. Connects to master on port 5001
 - b. Receives Map and Reduce tasks scheduled by master
 - c. Checks for availability of slots
 - d. Decreases the duration till the task is executed

e. Acquires and releases lock while decreasing duration

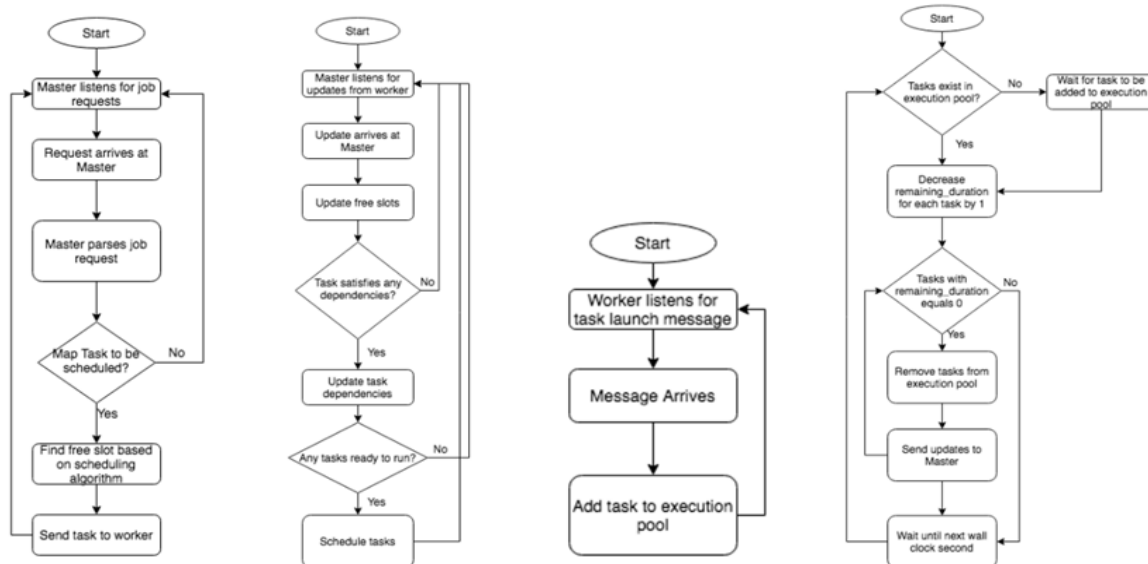
2. slot_handler :

- a. Assigns the slots to each task
- b. Keeps track of no of available slots and reports to the master
- c. Acquires and releases lock while assigning the slots

Scheduling Algorithms implemented:-



Model and WorkFlow:-



- The scheduler in the master has been introduced as the third thread in the implementation. This thread encompasses the scheduling algorithms

Execution

- The **workers** must initially be spawned by running **worker.py** with the appropriate port numbers and the slots.
- The **master.py** must be executed with the appropriate **config.json** whose ports and slots tally with the corresponding workers. The number of workers is not fixed to 3 and can be dynamic.
- Then **request.py** must be run with any amount of requests.
- The **Jobs** shall be executed and the corresponding log file will get created/overwritten based on the **algorithm**.
- Once all log files are created, the **analyzer.py** must be executed which reads all the log files and generates the results.

Results

1. Implementation of Round Robin Scheduler


a. Log File

	LL.log	RANDOM.log	RR.log
1	2020-12-01 19:43:28,907	Started_Job->	Job_Id: 0
2	2020-12-01 19:43:28,912	Started_Task->	Type: map Task_Id: 0_M0 Job_Id: 0 Worker_Id: 1
3	2020-12-01 19:43:28,912	Started_Task->	Type: map Task_Id: 0_M1 Job_Id: 0 Worker_Id: 2
4	2020-12-01 19:43:29,807	Started_Job->	Job_Id: 1
5	2020-12-01 19:43:29,813	Started_Task->	Type: map Task_Id: 1_M0 Job_Id: 1 Worker_Id: 3
6	2020-12-01 19:43:29,813	Started_Task->	Type: map Task_Id: 1_M1 Job_Id: 1 Worker_Id: 1
7	2020-12-01 19:43:29,949	Started_Job->	Job_Id: 2
8	2020-12-01 19:43:29,955	Started_Task->	Type: map Task_Id: 2_M0 Job_Id: 2 Worker_Id: 2
9	2020-12-01 19:43:29,955	Started_Task->	Type: map Task_Id: 2_M1 Job_Id: 2 Worker_Id: 3
10	2020-12-01 19:43:29,956	Started_Task->	Type: map Task_Id: 2_M2 Job_Id: 2 Worker_Id: 1
11	2020-12-01 19:43:30,311	Completed_Task->	Type: map Task_Id: 0_M0 Job_Id: 0 Worker_Id: 1
12	2020-12-01 19:43:30,695	Started_Job->	Job_Id: 3
13	2020-12-01 19:43:30,700	Started_Task->	Type: map Task_Id: 3_M0 Job_Id: 3 Worker_Id: 2
14	2020-12-01 19:43:30,872	Started_Job->	Job_Id: 4
15	2020-12-01 19:43:30,877	Started_Task->	Type: map Task_Id: 4_M0 Job_Id: 4 Worker_Id: 3
16	2020-12-01 19:43:30,878	Started_Task->	Type: map Task_Id: 4_M1 Job_Id: 4 Worker_Id: 1
17	2020-12-01 19:43:31,087	Completed_Task->	Type: map Task_Id: 0_M1 Job_Id: 0 Worker_Id: 2
18	2020-12-01 19:43:31,093	Started_Task->	Type: reduce Task_Id: 0_R0 Job_Id: 0 Worker_Id: 2
19	2020-12-01 19:43:31,307	Completed_Task->	Type: map Task_Id: 1_M1 Job_Id: 1 Worker_Id: 1
20	2020-12-01 19:43:32,238	Completed_Task->	Type: map Task_Id: 1_M0 Job_Id: 1 Worker_Id: 3
21	2020-12-01 19:43:32,244	Started_Task->	Type: reduce Task_Id: 0_R1 Job_Id: 0 Worker_Id: 3
22	2020-12-01 19:43:32,244	Started_Task->	Type: reduce Task_Id: 1_R0 Job_Id: 1 Worker_Id: 1
23	2020-12-01 19:43:33,098	Completed_Task->	Type: map Task_Id: 3_M0 Job_Id: 3 Worker_Id: 2
24	2020-12-01 19:43:33,104	Started_Task->	Type: reduce Task_Id: 1_R1 Job_Id: 1 Worker_Id: 2
25	2020-12-01 19:43:33,313	Completed_Task->	Type: map Task_Id: 2_M2 Job_Id: 2 Worker_Id: 1
26	2020-12-01 19:43:34,099	Completed_Task->	Type: map Task_Id: 2_M0 Job_Id: 2 Worker_Id: 2
27	2020-12-01 19:43:34,240	Completed_Task->	Type: map Task_Id: 2_M1 Job_Id: 2 Worker_Id: 3
28	2020-12-01 19:43:34,246	Completed_Task->	Type: map Task_Id: 4_M0 Job_Id: 4 Worker_Id: 3
29	2020-12-01 19:43:34,246	Started_Task->	Type: reduce Task_Id: 3_R0 Job_Id: 3 Worker_Id: 3
30	2020-12-01 19:43:34,246	Started_Task->	Type: reduce Task_Id: 3_R1 Job_Id: 3 Worker_Id: 1
31	2020-12-01 19:43:34,247	Started_Task->	Type: reduce Task_Id: 2_R0 Job_Id: 2 Worker_Id: 2

b. Mean and Median results :-

```
RR.log
Task mean : 3.1841578947368423
Task median : 3.369
Job mean : 7.819
Job median : 7.463
```

c. Plot of Scheduled tasks on Workers :



2. Implementation of Random Scheduler

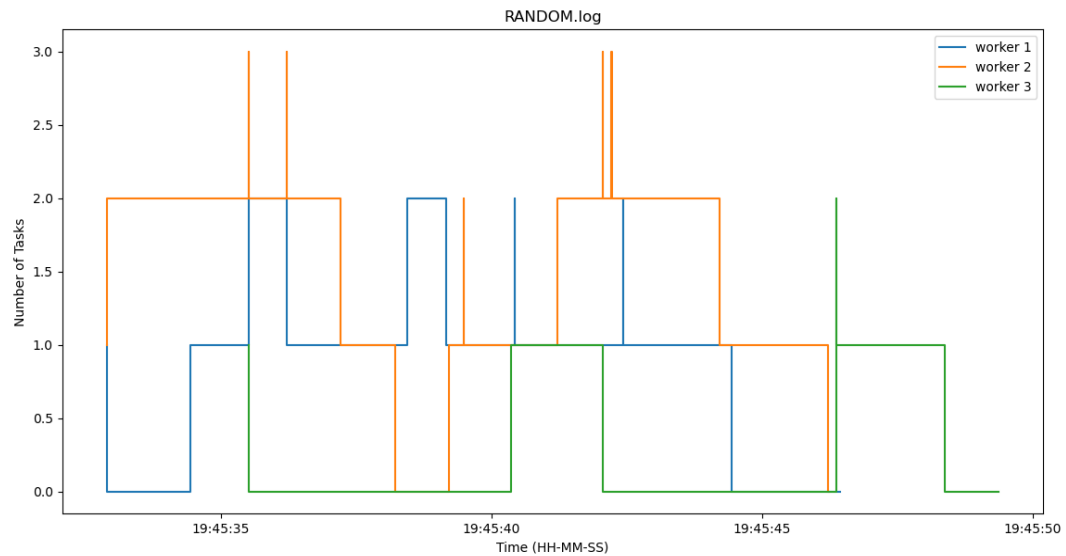
a. Log File

Line	Timestamp	Event	Job_Id	Task_Id	Job_Id	Worker_Id
1	2020-12-01 19:45:32,894	Started_Job->	Job_Id: 0			
2	2020-12-01 19:45:32,899	Started_Task->	Type: map	Task_Id: 0_M0	Job_Id: 0	Worker_Id: 1
3	2020-12-01 19:45:32,900	Started_Task->	Type: map	Task_Id: 0_M1	Job_Id: 0	Worker_Id: 2
4	2020-12-01 19:45:34,439	Completed_Task->	Type: map	Task_Id: 0_M0	Job_Id: 0	Worker_Id: 1
5	2020-12-01 19:45:35,515	Started_Job->	Job_Id: 1			
6	2020-12-01 19:45:35,520	Started_Task->	Type: map	Task_Id: 1_M0	Job_Id: 1	Worker_Id: 2
7	2020-12-01 19:45:35,521	Started_Task->	Type: map	Task_Id: 1_M1	Job_Id: 1	Worker_Id: 1
8	2020-12-01 19:45:35,521	Started_Task->	Type: map	Task_Id: 1_M2	Job_Id: 1	Worker_Id: 2
9	2020-12-01 19:45:35,522	Started_Task->	Type: map	Task_Id: 1_M3	Job_Id: 1	Worker_Id: 3
10	2020-12-01 19:45:36,218	Completed_Task->	Type: map	Task_Id: 0_M1	Job_Id: 0	Worker_Id: 2
11	2020-12-01 19:45:36,224	Started_Task->	Type: reduce	Task_Id: 0_R0	Job_Id: 0	Worker_Id: 2
12	2020-12-01 19:45:36,224	Started_Task->	Type: reduce	Task_Id: 0_R1	Job_Id: 0	Worker_Id: 1
13	2020-12-01 19:45:37,219	Completed_Task->	Type: map	Task_Id: 1_M2	Job_Id: 1	Worker_Id: 2
14	2020-12-01 19:45:38,221	Completed_Task->	Type: map	Task_Id: 1_M0	Job_Id: 1	Worker_Id: 2
15	2020-12-01 19:45:38,443	Completed_Task->	Type: reduce	Task_Id: 0_R1	Job_Id: 0	Worker_Id: 1
16	2020-12-01 19:45:39,162	Started_Job->	Job_Id: 2			
17	2020-12-01 19:45:39,168	Started_Task->	Type: map	Task_Id: 2_M0	Job_Id: 2	Worker_Id: 1
18	2020-12-01 19:45:39,221	Completed_Task->	Type: reduce	Task_Id: 0_R0	Job_Id: 0	Worker_Id: 2
19	2020-12-01 19:45:39,242	Completed_Job->	Job_Id: 0			
20	2020-12-01 19:45:39,495	Started_Job->	Job_Id: 3			
21	2020-12-01 19:45:39,500	Started_Task->	Type: map	Task_Id: 3_M0	Job_Id: 3	Worker_Id: 2
22	2020-12-01 19:45:39,501	Started_Task->	Type: map	Task_Id: 3_M1	Job_Id: 3	Worker_Id: 2
23	2020-12-01 19:45:40,370	Completed_Task->	Type: map	Task_Id: 1_M3	Job_Id: 1	Worker_Id: 3
24	2020-12-01 19:45:40,435	Completed_Task->	Type: map	Task_Id: 1_M1	Job_Id: 1	Worker_Id: 1
25	2020-12-01 19:45:40,441	Started_Task->	Type: reduce	Task_Id: 1_R0	Job_Id: 1	Worker_Id: 1
26	2020-12-01 19:45:41,224	Completed_Task->	Type: map	Task_Id: 3_M1	Job_Id: 3	Worker_Id: 2
27	2020-12-01 19:45:42,054	Started_Job->	Job_Id: 4			
28	2020-12-01 19:45:42,061	Started_Task->	Type: map	Task_Id: 4_M0	Job_Id: 4	Worker_Id: 2
29	2020-12-01 19:45:42,062	Started_Task->	Type: map	Task_Id: 4_M1	Job_Id: 4	Worker_Id: 2
30	2020-12-01 19:45:42,063	Started_Task->	Type: map	Task_Id: 4_M2	Job_Id: 4	Worker_Id: 3
31	2020-12-01 19:45:42,225	Completed_Task->	Type: map	Task_Id: 3_M0	Job_Id: 3	Worker_Id: 2

b. Mean and Median results :

```
RANDOM.log  
Task mean : 3.1372631578947368  
Task median : 2.997  
Job mean : 6.9404  
Job median : 7.3
```

c. Plot of Scheduled tasks on Workers



3. Implementation of Least Loaded Scheduler

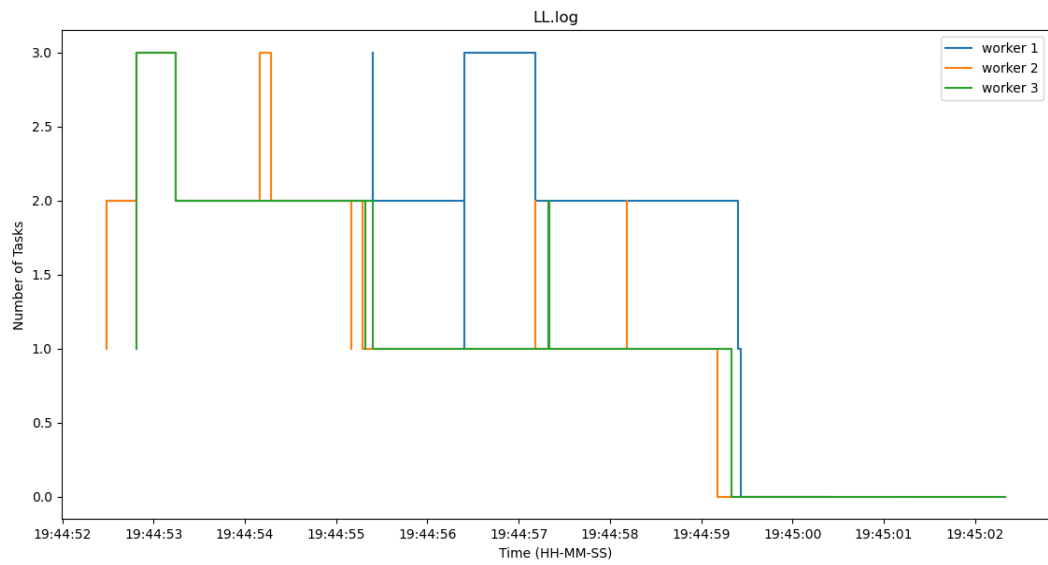
a. Log File

1	2020-12-01	19:44:52,478	Started_Job->	Job_Id: 0			
2	2020-12-01	19:44:52,484	Started_Task->	Type: map	Task_Id: 0_M0	Job_Id: 0	Worker_Id: 1
3	2020-12-01	19:44:52,484	Started_Task->	Type: map	Task_Id: 0_M1	Job_Id: 0	Worker_Id: 2
4	2020-12-01	19:44:52,811	Started_Job->	Job_Id: 1			
5	2020-12-01	19:44:52,816	Started_Task->	Type: map	Task_Id: 1_M0	Job_Id: 1	Worker_Id: 3
6	2020-12-01	19:44:52,817	Started_Task->	Type: map	Task_Id: 1_M1	Job_Id: 1	Worker_Id: 1
7	2020-12-01	19:44:52,817	Started_Task->	Type: map	Task_Id: 1_M2	Job_Id: 1	Worker_Id: 2
8	2020-12-01	19:44:52,817	Started_Task->	Type: map	Task_Id: 1_M3	Job_Id: 1	Worker_Id: 3
9	2020-12-01	19:44:53,234	Started_Job->	Job_Id: 2			
10	2020-12-01	19:44:53,239	Started_Task->	Type: map	Task_Id: 2_M0	Job_Id: 2	Worker_Id: 1
11	2020-12-01	19:44:53,240	Started_Task->	Type: map	Task_Id: 2_M1	Job_Id: 2	Worker_Id: 2
12	2020-12-01	19:44:53,240	Started_Task->	Type: map	Task_Id: 2_M2	Job_Id: 2	Worker_Id: 3
13	2020-12-01	19:44:53,279	Started_Job->	Job_Id: 3			
14	2020-12-01	19:44:53,811	Started_Job->	Job_Id: 4			
15	2020-12-01	19:44:54,162	Completed_Task->	Type: map	Task_Id: 0_M1	Job_Id: 0	Worker_Id: 2
16	2020-12-01	19:44:54,286	Started_Task->	Type: map	Task_Id: 3_M0	Job_Id: 3	Worker_Id: 2
17	2020-12-01	19:44:55,163	Completed_Task->	Type: map	Task_Id: 2_M1	Job_Id: 2	Worker_Id: 2
18	2020-12-01	19:44:55,163	Completed_Task->	Type: map	Task_Id: 1_M2	Job_Id: 1	Worker_Id: 2
19	2020-12-01	19:44:55,288	Started_Task->	Type: map	Task_Id: 4_M0	Job_Id: 4	Worker_Id: 2
20	2020-12-01	19:44:55,319	Completed_Task->	Type: map	Task_Id: 1_M3	Job_Id: 1	Worker_Id: 3
21	2020-12-01	19:44:55,319	Completed_Task->	Type: map	Task_Id: 2_M2	Job_Id: 2	Worker_Id: 3
22	2020-12-01	19:44:55,401	Completed_Task->	Type: map	Task_Id: 2_M0	Job_Id: 2	Worker_Id: 1
23	2020-12-01	19:44:55,406	Started_Task->	Type: reduce	Task_Id: 2_R0	Job_Id: 2	Worker_Id: 3
24	2020-12-01	19:44:55,406	Started_Task->	Type: reduce	Task_Id: 2_R1	Job_Id: 2	Worker_Id: 1
25	2020-12-01	19:44:56,402	Completed_Task->	Type: map	Task_Id: 0_M0	Job_Id: 0	Worker_Id: 1
26	2020-12-01	19:44:56,407	Completed_Task->	Type: map	Task_Id: 1_M1	Job_Id: 1	Worker_Id: 1
27	2020-12-01	19:44:56,407	Started_Task->	Type: reduce	Task_Id: 0_R0	Job_Id: 0	Worker_Id: 1
28	2020-12-01	19:44:57,180	Completed_Task->	Type: map	Task_Id: 4_M0	Job_Id: 4	Worker_Id: 2
29	2020-12-01	19:44:57,185	Started_Task->	Type: reduce	Task_Id: 4_R0	Job_Id: 4	Worker_Id: 2
30	2020-12-01	19:44:57,186	Started_Task->	Type: reduce	Task_Id: 4_R1	Job_Id: 4	Worker_Id: 1
31	2020-12-01	19:44:57,326	Completed_Task->	Type: map	Task_Id: 1_M0	Job_Id: 1	Worker_Id: 3

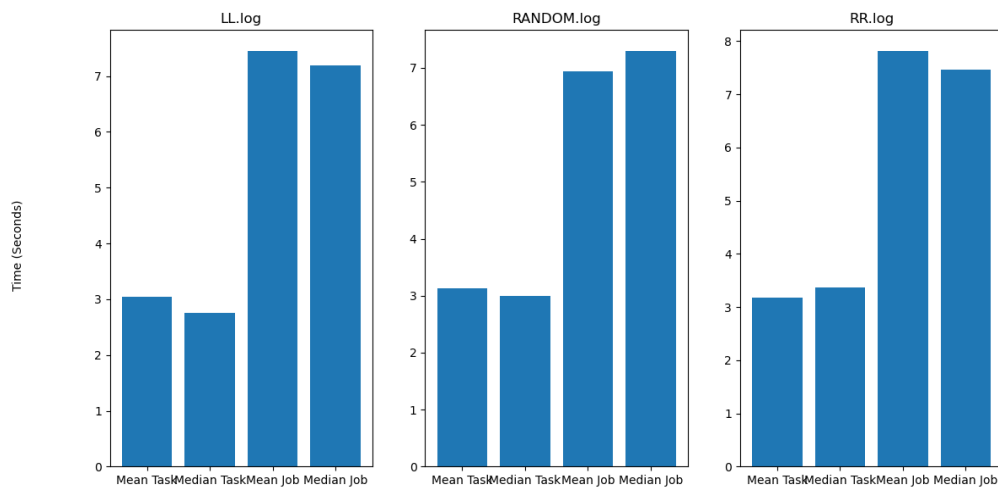
b. Mean and Median results :

```
LL.log
Task mean : 3.036777777777778
Task median : 2.75
Job mean : 7.451
Job median : 7.193
```

c. Plot of Scheduled tasks on Workers



Comparison of Performances of each Scheduling Algorithms



Problems

Problem	Solution
1. Standardized format of messages used to communicate between ports.	Json format was used for simplified encoding and decoding at sender and receiver
2. Parallelizing the implementation of multiple jobs	Creation and utilization of a 3rd thread on master to send tasks to workers
3. Consistency in format of logging messages	Using in-built logging module in python

Conclusion

Main Learnings (Observations) from the Project:-

- Resolving concurrency issues with the help of locking mechanism
- When the no. of slots in each worker is 1, the mean and median of the performance of each Scheduling algorithm is similar
- In some cases Random Scheduling algorithm proved to be more effective compared to Round Robin Scheduling
- In Round Robin, there was a fair distribution of load among all workers
- Least Loaded algorithm was found to be the most effective
- Random Scheduling algorithm caused an increase in total execution time in a few cases, due to queuing, increasing the waiting time unnecessarily in spite of workers being free.

EVALUATIONS:

SNo	Name	SRN	Contribution (Individual)
1	S Nikhil Ram	PES1201801972	Implementation of Master
2	Rakshith C	PES1201801903	Logging, Analyser and Scheduling algorithms
3	Vikas Gowda	PES1201801957	Implementation of Worker
4	Yash Gawankar	PES1201801482	Socket Programming and Resolving Concurrency using Locking

CHECKLIST:

SNo	Item	Status
1.	Source code documented	
2.	Source code uploaded to GitHub – (access link for the same, to be added in status)	
3.	Instructions for building and running the code. Your code must be usable out of the box.	